

# DPA Project

Group Project contributed by all

2023-04-22

```
#Install the pacman package
if(!require(pacman)) install.packages("pacman", repos = "http://cran.us.r-
project.org")

## Loading required package: pacman

## Warning: package 'pacman' was built under R version 4.2.3

#Load the required libraries
#If a package below is missing, p_load will automatically download it from
CRAN
pacman::p_load(tidyverse, ggplot2, ggthemes, data.table, lubridate, caret,
               knitr, scales, treemapify)
pacman::p_load(stringr)
pacman::p_load(dplyr)
```

## Data Preparation

```
#Download File
dl <- tempfile()
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)

#Construct a data frame called 'ratings' by utilizing the 'fread' function
from the data.table library
ratings <- fread(text = gsub("::", "\t", readLines(unzip(dl, "ml-
10M100K/ratings.dat"))),
                 col.names = c("userId", "movieId", "rating", "timestamp"))

movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")),
                          "\\:\\\\:", 3)

colnames(movies) <- c("movieId", "title", "genres")

movies <- as.data.frame(movies) %>%
  mutate(movieId = as.numeric(unique(movieId)),
         title = as.character(title),
         genres = as.character(genres))
movielens <- left_join(ratings, movies, by = "movieId")

#Designate the validation set as 10% of the MovieLens data
set.seed(1, sample.kind="Rounding")
```

```
## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding'
sampler
## used

test_index <-createDataPartition(y = movielens$rating, times = 1, p = 0.1,
list = FALSE)
edx <-movielens[ -test_index,]
temp <-movielens[test_index,]
#Check if userId and movieId in validation set are also in edx set
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

#Merge the rows that were removed from the validation set back into the edx
set
removed <-anti_join(temp, validation)

## Joining with `by = join_by(userId, movieId, rating, timestamp, title,
genres)`

edx <-rbind(edx, removed)
rm(dl, ratings, movies, test_index, temp, movielens, removed)

#divide Training and Test Sets:
set.seed(1, sample.kind = "Rounding")

## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding'
sampler
## used

test_index <-createDataPartition(y = edx$rating, times = 1, p = 0.1, list =
F)
edx_train <-edx[ -test_index,]
edx_temp <-edx[test_index,]
#Make sure userId and movieId are in the train and test sets
edx_test <-edx_temp %>%
  semi_join(edx_train, by = "movieId") %>%
  semi_join(edx_train, by = "userId")
removed <-anti_join(edx_temp, edx_test)

## Joining with `by = join_by(userId, movieId, rating, timestamp, title,
genres)`

edx_train <-rbind(edx_train, removed)
rm(edx_temp, test_index, removed)
```

## Analyzing the data

```
edx %>% as_tibble()

## # A tibble: 9,000,055 × 6
##   userId movieId rating timestamp title
```

```

genres
##      <int>   <dbl>  <dbl>      <int> <chr>
<chr>
##  1      1     122      5 838985046 Boomerang (1992)
Comed...
##  2      1     185      5 838983525 Net, The (1995)
Actio...
##  3      1     292      5 838983421 Outbreak (1995)
Actio...
##  4      1     316      5 838983392 Stargate (1994)
Actio...
##  5      1     329      5 838983392 Star Trek: Generations (1994)
Actio...
##  6      1     355      5 838984474 Flintstones, The (1994)
Child...
##  7      1     356      5 838983653 Forrest Gump (1994)
Comed...
##  8      1     362      5 838984885 Jungle Book, The (1994)
Adven...
##  9      1     364      5 838983707 Lion King, The (1994)
Adven...
## 10      1     370      5 838984596 Naked Gun 33 1/3: The Final Insult (1...
Actio...
## # ... with 9,000,045 more rows

```

*#Confirm the dimensions and explore the features and classes of edx.*  
glimpse(edx)

```

## Rows: 9,000,055
## Columns: 6
## $ userId      <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
2, ...
## $ movieId     <dbl> 122, 185, 292, 316, 329, 355, 356, 362, 364, 370, 377,
420, ...
## $ rating      <dbl> 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5, 5,
5, ...
## $ timestamp   <int> 838985046, 838983525, 838983421, 838983392, 838983392,
83898...
## $ title       <chr> "Boomerang (1992)", "Net, The (1995)", "Outbreak
(1995)", "S...
## $ genres      <chr> "Comedy|Romance", "Action|Crime|Thriller",
"Action|Drama|Sci...

```

*#Determine the unique number of userIds, movieIds, and genres*

```

edx %>% summarize(unique_users = length(unique(userId)),
                  unique_movies = length(unique(movieId)),
                  unique_genres = length(unique(genres)))

##   unique_users unique_movies unique_genres
## 1          69878          10677           797

```

```

#Ratings
length(unique(edx$rating))

## [1] 10

unique_ratings <- unique(edx$rating)
sort(unique_ratings)

## [1] 0.5 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0

#View a Tibble of the Ratings Distribution
edx %>% group_by(rating) %>% summarize(ratings_sum = n()) %>%
  arrange(desc(ratings_sum))

## # A tibble: 10 × 2
##   rating ratings_sum
##   <dbl>     <int>
## 1     4       2588430
## 2     3       2121240
## 3     5       1390114
## 4   3.5        791624
## 5     2        711422
## 6   4.5        526736
## 7     1        345679
## 8   2.5        333010
## 9   1.5        106426
## 10    0.5         85374

rp <- edx %>% filter(edx$rating >=3)
nrow(rp)/length(edx$rating)

## [1] 0.8242332

```

## TIMESTAMP

```

#Transform the timestamp column of the edx dataset to a 'RatingYear' format
edx <- edx %>% mutate(timestamp = as.POSIXct(timestamp, origin = "1970-01-
01",
                                     tz = "EST"))
edx$timestamp <- format(edx$timestamp, "%Y")
names(edx)[names(edx) == "timestamp"] <- "RatingYear"
head(edx)

##   userId movieId rating RatingYear      title
## 1:      1    122      5      1996 Boomerang (1992)
## 2:      1    185      5      1996 Net, The (1995)
## 3:      1    292      5      1996 Outbreak (1995)
## 4:      1    316      5      1996 Stargate (1994)
## 5:      1    329      5      1996 Star Trek: Generations (1994)
## 6:      1    355      5      1996 Flintstones, The (1994)
##
##               genres
## 1:      Comedy|Romance

```

```

## 2:      Action|Crime|Thriller
## 3: Action|Drama|Sci-Fi|Thriller
## 4:      Action|Adventure|Sci-Fi
## 5: Action|Adventure|Drama|Sci-Fi
## 6:      Children|Comedy|Fantasy

validation <- validation %>% mutate(timestamp = as.POSIXct(timestamp, origin
= "1970-01-01",
                                tz = "EST"))
validation$timestamp <- format(validation$timestamp, "%Y")
names(validation)[names(validation) == "timestamp"] <- "RatingYear"
head(validation)

##   userId movieId rating RatingYear
## 1:     1     231      5      1996
## 2:     1     480      5      1996
## 3:     1     586      5      1996
## 4:     2     151      3      1997
## 5:     2     858      2      1997
## 6:     2    1544      3      1997
##                                     title
## 1:                               Dumb & Dumber (1994)
## 2:                               Jurassic Park (1993)
## 3:                               Home Alone (1990)
## 4:                               Rob Roy (1995)
## 5:                               Godfather, The (1972)
## 6: Lost World: Jurassic Park, The (Jurassic Park 2) (1997)
##                                     genres
## 1:                               Comedy
## 2:      Action|Adventure|Sci-Fi|Thriller
## 3:                               Children|Comedy
## 4:      Action|Drama|Romance|War
## 5:                               Crime|Drama
## 6: Action|Adventure|Horror|Sci-Fi|Thriller

edx_train <- edx_train %>% mutate(timestamp = as.POSIXct(timestamp, origin =
"1970-01-01",
                                tz = "EST"))
edx_train$timestamp <- format(edx_train$timestamp, "%Y")
names(edx_train)[names(edx_train) == "timestamp"] <- "RatingYear"
head(edx_train)

##   userId movieId rating RatingYear
## 1:     1     122      5      1996
## 2:     1     292      5      1996
## 3:     1     316      5      1996
## 4:     1     329      5      1996
## 5:     1     355      5      1996
## 6:     1     356      5      1996
##                                     title
## 1:                               Boomerang (1992)
## 2:                               Outbreak (1995)
## 3:                               Stargate (1994)
## 4: Star Trek: Generations (1994)
## 5:                               Flintstones, The (1994)
## 6:                               Forrest Gump (1994)
##                                     genres
## 1:                               Comedy|Romance

```

```

## 2: Action|Drama|Sci-Fi|Thriller
## 3:      Action|Adventure|Sci-Fi
## 4: Action|Adventure|Drama|Sci-Fi
## 5:      Children|Comedy|Fantasy
## 6:      Comedy|Drama|Romance|War

edx_test <-edx_test %>% mutate(timestamp = as.POSIXct(timestamp, origin =
"1970-01-01",
                                tz = "EST"))
edx_test$timestamp <- format(edx_test$timestamp, "%Y")
names(edx_test)[names(edx_test) == "timestamp"] <- "RatingYear"
head(edx_test)

##      userId movieId rating RatingYear
## 1:      1      185      5      1996
## 2:      2      260      5      1997
## 3:      2      590      5      1997
## 4:      2     1049      3      1997
## 5:      2     1210      4      1997
## 6:      3     1148      4      2005
##                                     title
## 1:                                     Net, The (1995)
## 2: Star Wars: Episode IV - A New Hope (a.k.a. Star Wars) (1977)
## 3:                                     Dances with Wolves (1990)
## 4:                                     Ghost and the Darkness, The (1996)
## 5:      Star Wars: Episode VI - Return of the Jedi (1983)
## 6:      Wallace & Gromit: The Wrong Trousers (1993)
##                                     genres
## 1:      Action|Crime|Thriller
## 2:      Action|Adventure|Sci-Fi
## 3:      Adventure|Drama|Western
## 4:      Action|Adventure
## 5:      Action|Adventure|Sci-Fi
## 6: Animation|Children|Comedy|Crime

range(edx$RatingYear)

## [1] "1995" "2009"

#Convert the 'RatingYear' column from character to numeric data type in order to plot a histogram
edx$RatingYear <-as.numeric(edx$RatingYear)
str(edx)

## Classes 'data.table' and 'data.frame':  9000055 obs. of  6 variables:
## $ userId      : int  1 1 1 1 1 1 1 1 1 1 ...
## $ movieId      : num  122 185 292 316 329 355 356 362 364 370 ...
## $ rating       : num  5 5 5 5 5 5 5 5 5 5 ...
## $ RatingYear: num  1996 1996 1996 1996 1996 ...
## $ title       : chr  "Boomerang (1992)" "Net, The (1995)" "Outbreak (1995)"
## "Stargate (1994)" ...

```

```
## $ genres      : chr "Comedy|Romance" "Action|Crime|Thriller"
"Action|Drama|Sci-Fi|Thriller" "Action|Adventure|Sci-Fi" ...
## - attr(*, ".internal.selfref")=<externalptr>

edx %>% group_by(RatingYear, title) %>%
  summarize(Ratings_Sum = n(), Average_Rating = mean(rating)) %>%
  mutate(Average_Rating = sprintf("%.2f", Average_Rating)) %>%
  arrange(-Ratings_Sum) %>% print(n = 50)
```

## `summarise()` has grouped output by 'RatingYear'. You can override using the  
## `.groups` argument.

```
## # A tibble: 75,964 × 4
## # Groups:   RatingYear [15]
##   RatingYear title                                     Ratin...1
Avera...2
##           <dbl> <chr>                                     <int>
<chr>
## 1      1996 Batman (1989)                                     12016
3.26
## 2      1996 Dances with Wolves (1990)                       11524
3.79
## 3      1996 Apollo 13 (1995)                                   11393
3.99
## 4      1996 Pulp Fiction (1994)                               10925
4.01
## 5      1996 Fugitive, The (1993)                             10901
4.12
## 6      1996 True Lies (1994)                                  10838
3.57
## 7      1996 Forrest Gump (1994)                               9986
4.12
## 8      1996 Batman Forever (1995)                             9907
3.13
## 9      1996 Aladdin (1992)                                    9856
3.67
## 10     1996 Jurassic Park (1993)                             9771
3.84
## 11     1996 Ace Ventura: Pet Detective (1994)                9724
2.96
## 12     1996 Clear and Present Danger (1994)                 9484
3.71
## 13     1996 Die Hard: With a Vengeance (1995)               9467
3.48
## 14     1996 Silence of the Lambs, The (1991)                9341
4.29
## 15     1996 Beauty and the Beast (1991)                     8895
3.68
## 16     1996 Stargate (1994)                                  8845
```

3.33		
## 17	1996 Shawshank Redemption, The (1994)	8728
4.48		
## 18	1996 Outbreak (1995)	8386
3.56		
## 19	1996 Star Trek: Generations (1994)	8284
3.42		
## 20	1996 Cliffhanger (1993)	8172
3.21		
## 21	1996 Braveheart (1995)	8106
4.26		
## 22	1996 Firm, The (1993)	8097
3.54		
## 23	1996 Crimson Tide (1995)	8039
3.82		
## 24	1996 Terminator 2: Judgment Day (1991)	7994
3.96		
## 25	1996 Speed (1994)	7949
3.79		
## 26	1996 Dumb & Dumber (1994)	7938
2.83		
## 27	1996 Net, The (1995)	7902
3.34		
## 28	1996 Lion King, The (1994)	7692
3.81		
## 29	1996 While You Were Sleeping (1995)	7674
3.61		
## 30	1996 Waterworld (1995)	7601
3.07		
## 31	1996 Interview with the Vampire: The Vampire Chronicle...	7544
3.41		
## 32	1996 GoldenEye (1995)	7421
3.44		
## 33	1996 Mrs. Doubtfire (1993)	7391
3.62		
## 34	1996 Seven (a.k.a. Se7en) (1995)	7022
3.96		
## 35	1996 Pretty Woman (1990)	6998
3.47		
## 36	1996 Mask, The (1994)	6945
3.34		
## 37	1996 Ghost (1990)	6840
3.61		
## 38	1996 Natural Born Killers (1994)	6497
3.15		
## 39	1996 Quiz Show (1994)	6417
3.65		
## 40	1996 Babe (1995)	6363
3.87		
## 41	1996 Sleepless in Seattle (1993)	6334



```

3.70
## 42      1996 Addams Family Values (1993)      6072
3.06
## 43      1996 Schindler's List (1993)          5894
4.52
## 44      1996 Four Weddings and a Funeral (1994)  5871
3.70
## 45      1996 12 Monkeys (Twelve Monkeys) (1995)  5861
3.90
## 46      1996 Get Shorty (1995)                5817
3.67
## 47      1996 Usual Suspects, The (1995)        5669
4.30
## 48      1996 Home Alone (1990)                 5430
3.15
## 49      1996 Disclosure (1994)                 5373
3.37
## 50      1996 Clueless (1995)                   5360
3.44
## # ... with 75,914 more rows, and abbreviated variable names 1Ratings_Sum,
## #    2Average_Rating

```

```
edx_genres <-edx %>% separate_rows(genres, sep = "\\|")
```

Sum of Movie Ratings per Genre

```

edx_genres %>%
  group_by(genres) %>% summarize(Ratings_Sum = n(), Average_Rating =
mean(rating)) %>%
  arrange(-Ratings_Sum)

## # A tibble: 20 × 3
##   genres      Ratings_Sum Average_Rating
##   <chr>          <int>         <dbl>
## 1 Drama      3910127          3.67
## 2 Comedy     3540930          3.44
## 3 Action     2560545          3.42
## 4 Thriller   2325899          3.51
## 5 Adventure  1908892          3.49
## 6 Romance    1712100          3.55
## 7 Sci-Fi     1341183          3.40
## 8 Crime      1327715          3.67
## 9 Fantasy     925637          3.50
## 10 Children   737994          3.42
## 11 Horror     691485          3.27
## 12 Mystery    568332          3.68
## 13 War        511147          3.78
## 14 Animation  467168          3.60
## 15 Musical    433080          3.56
## 16 Western    189394          3.56
## 17 Film-Noir  118541          4.01


```

```
## 18 Documentary          93066          3.78
## 19 IMAX                 8181           3.77
## 20 (no genres listed)    7            3.64

library(treemap)

## Warning: package 'treemap' was built under R version 4.2.3

# sum of ratings by genre
genre_ratings <- aggregate(rating ~ genres, edx, sum)
# construct treemap
treemap(genre_ratings, index = "genres", vSize = "rating",
        type = "value", palette = "Set3",
        title = "Genre Ratings")
```



```
#Arrange the Genres by Mean Rating
edx_genres %>%
  group_by(genres) %>% summarize(Ratings_Sum = n(), Average_Rating =
mean(rating)) %>%
  arrange(-Average_Rating)

## # A tibble: 20 × 3
##   genres          Ratings_Sum Average_Rating
##   <chr>          <int>         <dbl>
## 1 Film-Noir      118541          4.01
## 2 Documentary    93066          3.78
## 3 War           511147          3.78
## 4 IMAX           8181           3.77
```

```
## 5 Mystery          568332          3.68
## 6 Drama            3910127         3.67
## 7 Crime            1327715         3.67
## 8 (no genres listed)      7         3.64
## 9 Animation        467168         3.60
## 10 Musical          433080         3.56
## 11 Western          189394         3.56
## 12 Romance          1712100        3.55
## 13 Thriller         2325899        3.51
## 14 Fantasy           925637        3.50
## 15 Adventure        1908892        3.49
## 16 Comedy           3540930        3.44
## 17 Action           2560545        3.42
## 18 Children          737994        3.42
## 19 Sci-Fi           1341183        3.40
## 20 Horror            691485        3.27
```

*#Coerce the 'genres' column from character data type to factor data type*

```
edx$genres <- as.factor(edx$genres)
edx_genres$genres <- as.factor(edx_genres$genres)
class(edx_genres$genres)
```


```
## [1] "factor"
```

```
library(ggplot2)
```

*# Aggregate of ratings per genre*

```
genre_ratings <- edx %>%
  separate_rows(genres, sep = "\\|") %>%
  group_by(genres) %>%
  summarize(total_ratings = sum(rating))
```

```
ggplot(genre_ratings, aes(x = reorder(genres, -total_ratings), y =
total_ratings, fill = genres)) +
  geom_bar(stat = "identity") +
  ggtitle("Sum of Movie Ratings per Genre") +
  xlab("Genre") +
  ylab("Total Ratings") +
  theme(plot.title = element_text(hjust = 0.5))
```




Mean Rating per Genre

```
library(ggplot2)

mean_ratings <- edx %>%
  separate_rows(genres, sep = "\\|") %>%
  group_by(genres) %>%
  summarize(mean_rating = mean(rating), .groups = 'drop')

ggplot(mean_ratings, aes(x = reorder(genres, mean_rating), y = mean_rating,
  fill = genres)) +
  geom_bar(stat = 'identity') +
  coord_flip() +
  ggtitle("Mean Rating per Genre") +
  xlab("Genre") +
  ylab("Mean Rating") +
  theme(plot.title = element_text(hjust = 0.5))
```



```
yearreleaseda <-as.numeric(str_sub(edx$title, start = -5, end = -2))
edx <- edx %>% mutate(yearReleased = yearreleaseda)
head(edx)
```

```
##      userId movieId rating RatingYear      title
## 1:      1      122      5      1996 Boomerang (1992)
## 2:      1      185      5      1996 Net, The (1995)
## 3:      1      292      5      1996 Outbreak (1995)
## 4:      1      316      5      1996 Stargate (1994)
## 5:      1      329      5      1996 Star Trek: Generations (1994)
## 6:      1      355      5      1996 Flintstones, The (1994)
##      genres yearReleased
## 1:      Comedy|Romance      1992
## 2:      Action|Crime|Thriller      1995
## 3: Action|Drama|Sci-Fi|Thriller      1995
## 4:      Action|Adventure|Sci-Fi      1994
## 5: Action|Adventure|Drama|Sci-Fi      1994
## 6:      Children|Comedy|Fantasy      1994
```

*#Do the same for the validation set*

```
yearreleasedb <-as.numeric(str_sub(validation$title, start = -5, end = -2))
validation <- validation %>% mutate(yearReleased = yearreleasedb)
head(validation)
```

```
##      userId movieId rating RatingYear
## 1:      1      231      5      1996
## 2:      1      480      5      1996
```

```

## 3:      1      586      5      1996
## 4:      2      151      3      1997
## 5:      2      858      2      1997
## 6:      2     1544      3      1997
##
##                                     title
## 1:                                Dumb & Dumber (1994)
## 2:                                Jurassic Park (1993)
## 3:                                Home Alone (1990)
## 4:                                Rob Roy (1995)
## 5:                                Godfather, The (1972)
## 6: Lost World: Jurassic Park, The (Jurassic Park 2) (1997)
##                                     genres yearReleased
## 1:                                Comedy          1994
## 2:      Action|Adventure|Sci-Fi|Thriller          1993
## 3:                                Children|Comedy          1990
## 4:      Action|Drama|Romance|War          1995
## 5:                                Crime|Drama          1972
## 6: Action|Adventure|Horror|Sci-Fi|Thriller          1997

#This is also applied to edx_train & edx_test for later modeling purposes
yearreleasedc <-as.numeric(str_sub(edx_train$title, start = -5, end = -2))
edx_train <- edx_train %>% mutate(yearReleased = yearreleasedc)
head(edx_train)

##      userId movieId rating RatingYear                                     title
## 1:      1      122      5      1996                                Boomerang (1992)
## 2:      1      292      5      1996                                Outbreak (1995)
## 3:      1      316      5      1996                                Stargate (1994)
## 4:      1      329      5      1996 Star Trek: Generations (1994)
## 5:      1      355      5      1996      Flintstones, The (1994)
## 6:      1      356      5      1996      Forrest Gump (1994)
##                                     genres yearReleased
## 1:                                Comedy|Romance          1992
## 2:      Action|Drama|Sci-Fi|Thriller          1995
## 3:      Action|Adventure|Sci-Fi          1994
## 4: Action|Adventure|Drama|Sci-Fi          1994
## 5:      Children|Comedy|Fantasy          1994
## 6:      Comedy|Drama|Romance|War          1994

yearreleasedd <-as.numeric(str_sub(edx_test$title, start = -5, end = -2))
edx_test <- edx_test %>% mutate(yearReleased = yearreleasedd)
head(edx_test)

##      userId movieId rating RatingYear
## 1:      1      185      5      1996
## 2:      2      260      5      1997
## 3:      2      590      5      1997
## 4:      2     1049      3      1997
## 5:      2     1210      4      1997
## 6:      3     1148      4      2005
##
##                                     title

```

```
## 1: Net, The (1995)
## 2: Star Wars: Episode IV - A New Hope (a.k.a. Star Wars) (1977)
## 3: Dances with Wolves (1990)
## 4: Ghost and the Darkness, The (1996)
## 5: Star Wars: Episode VI - Return of the Jedi (1983)
## 6: Wallace & Gromit: The Wrong Trousers (1993)
##
##          genres yearReleased
## 1: Action|Crime|Thriller      1995
## 2: Action|Adventure|Sci-Fi    1977
## 3: Adventure|Drama|Western     1990
## 4: Action|Adventure           1996
## 5: Action|Adventure|Sci-Fi    1983
## 6: Animation|Children|Comedy|Crime 1993
```

Use the newly defined “yearReleased” column to add a “MovieAge” column

```
edx <-edx %>% mutate(MovieAge = 2020 - yearReleased)
validation <-validation %>% mutate(MovieAge = 2020 - yearReleased)
edx_train <-edx_train %>% mutate(MovieAge = 2020 - yearReleased)
edx_test <-edx_test %>% mutate(MovieAge = 2020 - yearReleased)
```

Movie Age

```
summary(edx$MovieAge)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 12.00   22.00   26.00   29.78   33.00   105.00
```

Modeling The formula for RMSE can be defined as follows with  $\hat{y}_{u,i}$  the prediction of movie  $i$  by user  $u$ , and  $y_{u,i}$  the rating of movie  $i$ , by user  $u$ .  $N$  is then defined as the number of user/movie combinations and the sum of these different combinations.

```
RMSE <- function(true_ratings, predicted_ratings){
  sqrt(mean((true_ratings - predicted_ratings)^2))
}
```

Begin Modeling: Benchmarking Model

```
edx_train_mu <-mean(edx_train$rating)
NRMSE_M1 <- RMSE(edx_test$rating, edx_train_mu)
#Table the Results
results_table <-tibble(Model_Type = "NRMSE", RMSE = NRMSE_M1) %>%
  mutate(RMSE = sprintf("%.4f", RMSE))
results_table

## # A tibble: 1 × 2
##   Model_Type RMSE
##   <chr>      <chr>
## 1 NRMSE      1.0601
```

Median Table

```

edx_train_median <- median(edx_train$rating)
MM_M2 <- RMSE(edx_test$rating, edx_train_median)
#Table the Results
results_table <- tibble(Model_Type = c("NRMSE", "Median_Model"),
                        RMSE = c(NRMSE_M1, MM_M2)) %>%
  mutate(RMSE = sprintf("%0.4f", RMSE))

results_table

## # A tibble: 2 × 2
##   Model_Type    RMSE
##   <chr>        <chr>
## 1 NRMSE        1.0601
## 2 Median_Model 1.1668

```

### Movie Effects Model

```

bi <- edx_train %>% group_by(movieId) %>%
  summarize(b_i = mean(rating - edx_train_mu))

```

create the prediction

```

prediction_bi <- edx_train_mu + edx_test %>%
  left_join(bi, by = "movieId") %>% .$b_i
MEM_M3 <- RMSE(edx_test$rating, prediction_bi)
#Table the Results
results_table <- tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects"),
                        RMSE = c(NRMSE_M1, MM_M2, MEM_M3)) %>%
  mutate(RMSE = sprintf("%0.4f", RMSE))

results_table

## # A tibble: 3 × 2
##   Model_Type    RMSE
##   <chr>        <chr>
## 1 NRMSE        1.0601
## 2 Median_Model 1.1668
## 3 Movie Effects 0.9430

```

### Adding User Effects to the Movie Effects Model:

```

bu <- edx_train %>% left_join(bi, by = "movieId") %>% group_by(userId) %>%
  summarize(b_u = mean(rating - edx_train_mu - b_i))

```

Create the Prediction Then check the prediction against the test set to determine the RMSE and table the results.

```

prediction_bu <- edx_test %>% left_join(bi, by = "movieId") %>%
  left_join(bu, by = "userId") %>%
  mutate(predictions = edx_train_mu + b_i + b_u) %>% .$predictions
UEM_M4 <- RMSE(edx_test$rating, prediction_bu)
#Table the Results
results_table <- tibble(Model_Type = c("NRMSE", "Median_Model", "Movie

```



```
Effects", "Movie & User Effects"),
      RMSE = c(NRMSE_M1, MM_M2, MEM_M3, UEM_M4)) %>%
      mutate(RMSE = sprintf("%.4f", RMSE))

results_table

## # A tibble: 4 × 2
##   Model_Type      RMSE
##   <chr>          <chr>
## 1 NRMSE          1.0601
## 2 Median_Model   1.1668
## 3 Movie Effects   0.9430
## 4 Movie & User Effects 0.8647
```

Adding Movie Age Effects: (Movie, User & Movie Age Effects Model)

```
ba <- edx_train %>%
  left_join(bi, by="movieId") %>% left_join(bu, by="userId") %>%
  group_by(MovieAge) %>% summarize(b_a = mean(rating - b_i - b_u -
edx_train_mu))
```

Create the Prediction Check the prediction against the test set to determine the RMSE and table the results.

```
predictions_ma <- edx_test %>%
  left_join(bi, by = "movieId") %>% left_join(bu, by = "userId") %>%
  left_join(ba, by = "MovieAge") %>% mutate(predictions = edx_train_mu + b_i
+ b_u + b_a) %>%
  .$predictions
UMMAE_M5 <-RMSE(edx_test$rating, predictions_ma)
#Table the results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects",
      "Movie & User Effects",
      "User, Movie & Movie Age Effects"),
      RMSE = c(NRMSE_M1, MM_M2, MEM_M3, UEM_M4, UMMAE_M5))
%>%
      mutate(RMSE = sprintf("%.4f", RMSE))

results_table

## # A tibble: 5 × 2
##   Model_Type      RMSE
##   <chr>          <chr>
## 1 NRMSE          1.0601
## 2 Median_Model   1.1668
## 3 Movie Effects   0.9430
## 4 Movie & User Effects 0.8647
## 5 User, Movie & Movie Age Effects 0.8643
```

Movie & User Effects Model with Regularization:

```

lambdasR <- seq(0, 10, 1)
RMSES <- sapply(lambdasR, function(l){
  edx_train_mu <- mean(edx_train$rating)

  b_i <- edx_train %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - edx_train_mu)/(n() + 1))

  b_u <- edx_train %>%
    left_join(b_i, by='movieId') %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - edx_train_mu)/(n() + 1))

  predicted_ratings <- edx_test %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    mutate(pred = edx_train_mu + b_i + b_u) %>% .$pred

  return(RMSE(predicted_ratings, edx_test$rating))
})
#Determine which Lambda minimizes the RMSE
lambda <- lambdasR[which.min(RMSES)]
lambda


## [1] 5

library(ggplot2)

# Create a data frame with LambdasR and RMSES
data <- data.frame(lambdasR = lambdasR, RMSES = RMSES)

# Create the scatter plot
ggplot(data, aes(x = lambdasR, y = RMSES)) +
  geom_point(color = "purple", alpha = 0.3) +
  ggtitle("RMSE vs. Lambda") +
  xlab("Lambda") + ylab("RMSE")

```



### Building the Movie & User Effects Model with Regularization

```

b_i <- edx_train %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - edx_train_mu)/(n()+lambda))
b_u <-edx_train %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - edx_train_mu)/(n()+lambda))
reg_prediction <- edx_test %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  mutate(predictions = edx_train_mu + b_i + b_u) %>% .$predictions

UMEM_REG_M6 <-RMSE(edx_test$rating, reg_prediction)
#Table the Results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects",
                                     "Movie & User Effects",
                                     "Movie, User & Movie Age Effects",
                                     "Movie & User Effects
w/Regularization"),
                      RMSE = c(NRMSE_M1, MM_M2, MEM_M3, UEM_M4,
                                UMMAE_M5, UMEM_REG_M6)) %>%
  mutate(RMSE = sprintf("%.6f", RMSE))
results_table

```

```
## # A tibble: 6 × 2
##   Model_Type          RMSE
##   <chr>             <chr>
## 1 NRMSE             1.060054
## 2 Median_Model      1.166756
## 3 Movie Effects      0.942961
## 4 Movie & User Effects 0.864684
## 5 Movie, User & Movie Age Effects 0.864330
## 6 Movie & User Effects w/Regularization 0.864136
```

Movie, User & Movie Age Effects Model with Regularization:

```
lambdasM <-seq(0, 10, 1)
RMSES2 <-sapply(lambdasM, function(l){
  edx_train_mu <-mean(edx_train$rating)

  b_i <-edx_train %>%
    group_by(movieId) %>%
    summarize(b_i = sum(rating - edx_train_mu)/(n() + 1))

  b_u <-edx_train %>%
    left_join(b_i, by='movieId') %>%
    group_by(userId) %>%
    summarize(b_u = sum(rating - b_i - edx_train_mu)/(n() +1))

  b_a <-edx_train %>%
    left_join(b_i, by = "movieId") %>% left_join(b_u, by = "userId") %>%
    group_by(MovieAge) %>%
    summarize(b_a = sum(rating - b_i - b_u - edx_train_mu)/(n()+1))

  predicted_ratings <-edx_test %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    left_join(b_a, by = "MovieAge") %>%
    mutate(predictions = edx_train_mu + b_i + b_u + b_a) %>% .$predictions

  return(RMSE(predicted_ratings, edx_test$rating))
})
lambda2 <- lambdasM[which.min(RMSES2)]
lambda2


## [1] 5

library(ggplot2)

# Create a data frame with LambdasM and RMSES2
df <- data.frame(lambda = lambdasM, RMSE = RMSES2)

# Create a scatterplot of RMSEs vs LambdasM
ggplot(df, aes(x = lambda, y = RMSE)) +
```

```
geom_point() +
geom_vline(xintercept = lambda2, linetype = "dashed") +
ggtitle("RMSE vs lambda") +
xlab("Lambda") + ylab("RMSE")
```



Building the User, Movie & Movie Age Effects Model with Regularization

```
b_i <- edx_train %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - edx_train_mu)/(n()+lambda2))
b_u <-edx_train %>%
  left_join(b_i, by = "movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - edx_train_mu)/(n()+lambda2))
b_a <-edx_train %>%
  left_join(b_i, by = "movieId") %>% left_join(b_u, by = "userId") %>%
  group_by(MovieAge) %>%
  summarize(b_a = sum(rating - b_i - b_u - edx_train_mu)/(n()+lambda2))
reg_prediction2 <- edx_test %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_a, by = "MovieAge") %>%
  mutate(pred = edx_train_mu + b_i + b_u + b_a) %>%
  pull(pred)
UMMAE_REG_M7 <-RMSE(edx_test$rating, reg_prediction2)
#Table the Results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects",
```

```

    "Movie & User Effects",
    "User, Movie & Movie Age Effects",
    "Movie & User Effects
w/Regularization",
    "User, Movie & Movie Age Effects
w/Regularization"),
  RMSE = c(NRMSE_M1, MM_M2, MEM_M3, UEM_M4,
    UMAE_M5, UMEM_REG_M6, UMAE_REG_M7)) %>%
  mutate(RMSE = sprintf("%0.5f", RMSE))
results_table

## # A tibble: 7 × 2
##   Model_Type      RMSE
##   <chr>          <chr>
## 1 NRMSE          1.06005
## 2 Median_Model   1.16676
## 3 Movie Effects   0.94296
## 4 Movie & User Effects 0.86468
## 5 User, Movie & Movie Age Effects 0.86433
## 6 Movie & User Effects w/Regularization 0.86414
## 7 User, Movie & Movie Age Effects w/Regularization 0.86384

```

Using Validation: Now we will move on to using the edx & validation sets to confirm our Final Model achieves an RMSE less than .8649.

The Benchmarking Model with Validation:

```

edx_mu <-mean(edx$rating)
FRMSE_M1 <-RMSE(validation$rating, edx_mu)
#Table the Results
results_table <-tibble(Model_Type = ("NRMSE"),
  Final_RMSE_Validation = (NRMSE_M1)) %>%
  mutate(Final_RMSE_Validation = sprintf("%0.5f",
Final_RMSE_Validation))
results_table

## # A tibble: 1 × 2
##   Model_Type Final_RMSE_Validation
##   <chr>          <chr>
## 1 NRMSE          1.06005

```

Median Model with validation:

```

edx_med <-median(edx$rating)
FRMSE_M2 <-RMSE(validation$rating, edx_med)
#Table the Results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model"),
  Final_RMSE_Validation = c(FRMSE_M1, FRMSE_M2)) %>%
  mutate(Final_RMSE_Validation = sprintf("%0.5f", Final_RMSE_Validation))
results_table

```

```
## # A tibble: 2 × 2
##   Model_Type   Final_RMSE_Validation
##   <chr>        <chr>
## 1 NRMSE        1.06120
## 2 Median_Model 1.16802

bi <- edx %>% group_by(movieId) %>%
  summarize(b_i = mean(rating - edx_mu))
#Prediction
prediction_bi <-edx_mu + validation %>%
  left_join(bi, by = "movieId") %>% .$b_i
FRMSE_M3 <-RMSE(validation$rating, prediction_bi)
#Table the Results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects"),
                        Final_RMSE_Validation = c(FRMSE_M1, FRMSE_M2,
FRMSE_M3)) %>%
  mutate(Final_RMSE_Validation = sprintf("%0.5f", Final_RMSE_Validation))
results_table

## # A tibble: 3 × 2
##   Model_Type   Final_RMSE_Validation
##   <chr>        <chr>
## 1 NRMSE        1.06120
## 2 Median_Model 1.16802
## 3 Movie Effects 0.94391
```

### Movie & User Effects Model with Validation

```
bu <-edx %>% left_join(bi, by = "movieId") %>% group_by(userId) %>%
  summarize(b_u = mean(rating - edx_mu - b_i))
#Prediction
prediction_bu <-validation %>% left_join(bi, by = "movieId") %>%
  left_join(bu, by = "userId") %>%
  mutate(predictions = edx_mu + b_i + b_u) %>% .$predictions
FRMSE_M4 <-RMSE(validation$rating, prediction_bu)
#Table the Results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects",
                                     "Movie & User Effects"),
                        Final_RMSE_Validation = c(FRMSE_M1, FRMSE_M2,
FRMSE_M3,
                                     FRMSE_M4)) %>%
  mutate(Final_RMSE_Validation = sprintf("%0.5f", Final_RMSE_Validation))
results_table

## # A tibble: 4 × 2
##   Model_Type   Final_RMSE_Validation
##   <chr>        <chr>
## 1 NRMSE        1.06120
## 2 Median_Model 1.16802
```

```
## 3 Movie Effects          0.94391
## 4 Movie & User Effects 0.86535
```

Movie, User & Movie Age Effects with Validation:

```
ba <- edx %>%
  left_join(bi, by = "movieId") %>% left_join(bu, by = "userId") %>%
  group_by(MovieAge) %>% summarize(b_a = mean(rating - b_i - b_u - edx_mu))
#Prediction
predictions_ma <- validation %>%
  left_join(bi, by = "movieId") %>% left_join(bu, by = "userId") %>%
  left_join(ba, by = "MovieAge") %>% mutate(predictions = edx_mu + b_i + b_u
+ b_a) %>%
  .$predictions
FRMSE_M5 <-RMSE(validation$rating, predictions_ma)
#Table the Results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects",
                                     "Movie & User Effects",
                                     "Movie, User, & Movie Age Effects"),
                       Final_RMSE_Validation = c(FRMSE_M1, FRMSE_M2,
FRMSE_M3,
                                                FRMSE_M4, FRMSE_M5)) %>%
  mutate(Final_RMSE_Validation = sprintf("%0.5f", Final_RMSE_Validation))
results_table

## # A tibble: 5 × 2
##   Model_Type                               Final_RMSE_Validation
##   <chr>                                     <chr>
## 1 NRMSE                                   1.06120
## 2 Median_Model                         1.16802
## 3 Movie Effects                        0.94391
## 4 Movie & User Effects                 0.86535
## 5 Movie, User, & Movie Age Effects 0.86500
```

Movie & User Effects with Regularization (Validation):

```
lambda
## [1] 5
#Movie & User Effects Model with Regularization using the validation set
b_i <-edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - edx_mu)/(n()+lambda))
b_u <-edx %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - edx_mu)/(n()+lambda))
reg_prediction <-validation %>%
```



```

left_join(b_i, by = "movieId") %>%
left_join(b_u, by = "userId") %>%
mutate(predictions = edx_mu + b_i + b_u) %>% .$predictions

FRMSE_M6 <-RMSE(validation$rating, reg_prediction)
#Table the Results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects",
                                     "Movie & User Effects",
                                     "Movie, User, & Movie Age Effects",
                                     "Movie & User Effects
w/Regularization"),
                       Final_RMSE_Validation = c(FRMSE_M1, FRMSE_M2,
FRMSE_M3,
                                                FRMSE_M4, FRMSE_M5,
                                                FRMSE_M6)) %>%
mutate(Final_RMSE_Validation = sprintf("%0.5f",
                                       Final_RMSE_Validation))

results_table

## # A tibble: 6 × 2
##   Model_Type                               Final_RMSE_Validation
##   <chr>                                <chr>
## 1 NRMSE                               1.06120
## 2 Median_Model                       1.16802
## 3 Movie Effects                      0.94391
## 4 Movie & User Effects                0.86535
## 5 Movie, User, & Movie Age Effects    0.86500
## 6 Movie & User Effects w/Regularization 0.86482

```

Final Model with Validation: This Model features Movie, User, & Movie Age Effects with Regularization

```

lambda2

## [1] 5

b_i <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - edx_mu)/(n()+lambda2))
b_u <-edx %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - edx_mu)/(n()+lambda2))
b_a <-edx %>%
  left_join(b_i, by="movieId") %>% left_join(b_u, by= "userId") %>%
  group_by(MovieAge) %>%
  summarize(b_a = sum(rating - b_i - b_u - edx_mu)/(n()+lambda2))
reg_prediction2 <-validation %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%

```

```

left_join(b_a, by = "MovieAge") %>%
mutate(predictions = edx_mu + b_i + b_u + b_a) %>% .$predictions

FRMSE_M7 <-RMSE(validation$rating, reg_prediction2)
#Table the Results
results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects",
                                     "Movie & User Effects",
                                     "Movie, User, & Movie Age Effects",
                                     "Movie & User Effects
w/Regularization",
                                     "Movie, User & Movie Age Effects
w/Regularization"),
                       Final_RMSE_Validation = c(FRMSE_M1, FRMSE_M2,
FRMSE_M3,
                                                FRMSE_M4, FRMSE_M5,
                                                FRMSE_M6, FRMSE_M7)) %>%
mutate(Final_RMSE_Validation = sprintf("%.5f",
                                       Final_RMSE_Validation))
results_table
## # A tibble: 7 × 2
##   Model_Type                               Final_RMSE_Validation
##   <chr>                                <chr>
## 1 NRMSE                               1.06120
## 2 Median_Model                       1.16802
## 3 Movie Effects                      0.94391
## 4 Movie & User Effects                0.86535
## 5 Movie, User, & Movie Age Effects    0.86500
## 6 Movie & User Effects w/Regularization 0.86482
## 7 Movie, User & Movie Age Effects w/Regularization 0.86452

#Building the User, Movie & Movie Age Effects Model with Regularization

```

Table the training & test set results against those of the validation set

```

results_table <-tibble(Model_Type = c("NRMSE", "Median_Model", "Movie
Effects",
                                     "Movie & User Effects",
                                     "Movie, User & Movie Age Effects",
                                     "Movie & User Effects
w/Regularization",
                                     "User, Movie & Movie Age Effects
w/Regularization"),
                       RMSE = c(NRMSE_M1, MM_M2, MEM_M3, UEM_M4,
                                UMMAE_M5, UMEM_REG_M6, UMMAE_REG_M7),
                       Final_RMSE_Validation = c(FRMSE_M1, FRMSE_M2,
FRMSE_M3, FRMSE_M4,
                                                FRMSE_M5, FRMSE_M6,
                                                FRMSE_M7)) %>%
mutate(Final_RMSE_Validation = sprintf("%.5f",

```

```

mutate(RMSE = sprintf("%0.5f", RMSE))

Final_RMSE_Validation)) %>%

results_table

## # A tibble: 7 × 3
##   Model_Type                RMSE
##   <chr>                  <chr>  <chr>
## 1 NRMSE                  1.06005 1.06120
## 2 Median_Model          1.16676 1.16802
## 3 Movie Effects          0.94296 0.94391
## 4 Movie & User Effects    0.86468 0.86535
## 5 Movie, User & Movie Age Effects 0.86433 0.86500
## 6 Movie & User Effects w/Regularization 0.86414 0.86482
## 7 User, Movie & Movie Age Effects w/Regularization 0.86384 0.86452

#The kable function in knitr table of the final results
results_table %>% knitr::kable()

```

Model_Type	RMSE	Final_RMSE_Validation
NRMSE	1.06005	1.06120
Median_Model	1.16676	1.16802
Movie Effects	0.94296	0.94391
Movie & User Effects	0.86468	0.86535
Movie, User & Movie Age Effects	0.86433	0.86500
Movie & User Effects w/Regularization	0.86414	0.86482
User, Movie & Movie Age Effects w/Regularization	0.86384	0.86452

```

# Create a data frame with the given data
model_data <- data.frame(
  Model_Type = c("NRMSE", "Median_Model", "Movie Effects", "Movie & User
Effects",
                "Movie, User & Movie Age Effects", "Movie & User Effects
w/Regularization",
                "User, Movie & Movie Age Effects w/Regularization"),
  RMSE = c(1.06005, 1.16676, 0.94296, 0.86468, 0.86433, 0.86414, 0.86384),
  Final_RMSE_Validation = c(1.06120, 1.16802, 0.94391, 0.86535, 0.86500,
0.86482, 0.86452)
)


# Create a bar plot of RMSE
barplot(
  model_data$RMSE,
  names.arg = model_data$Model_Type,
  xlab = "Model Type",
  ylab = "RMSE",

```

```

main = "RMSE by Model Type",
col = "blue"
)


```



```

# Create a bar plot of Final_RMSE_Validation
barplot(
  model_data$Final_RMSE_Validation,
  names.arg = model_data$Model_Type,
  xlab = "Model Type",
  ylab = "Final RMSE Validation",
  main = "Final RMSE Validation by Model Type",
  col = "red"
)

```



This Final Model achieves an RMSE of .86452 The lowest RMSE using the validation set is the Final Validation Model featuring Regularized User, Movie & Movie Age Effects.